

# Lecture 5-3

## Pandas: Indexing, Arithmetic, Missing Values

Week 5 Friday

Miles Chen, PhD

Based on Wes McKinney's Python for Data Analysis and the Pandas Documentation

In [1]:

```
import numpy as np
import pandas as pd
```

## Series that we will use as examples

In [2]:

```
# note that the value after the decimal place corresponds to the letter position.  
# i.e. 1.4 corresponds to d, the fourth letter.  
original1 = pd.Series([1.4, 2.3, 3.1, 4.2], index = ['d','c','a','b'])  
original2 = pd.Series([2.2, 3.1, 1.3, 4.4], index = ['b','a','c','d'])
```

In [3]:

```
original1 # when you create a series, the original order of the index is preserved
```

Out[3]:

```
d      1.4  
c      2.3  
a      3.1  
b      4.2  
dtype: float64
```

In [4]:

```
# making a DataFrame with multiple series with the same index preserves the index order  
pd.DataFrame({"x":original1, "x2": original1 * 2})
```

Out[4]:

x	x2
---	----

	x	x <sup>2</sup>
d	1.4	2.8
c	2.3	4.6
a	3.1	6.2
b	4.2	8.4

In [5]:

```
original2 # note that original1 and original2 have different index orders
```

Out[5]:

```
b      2.2
a      3.1
c      1.3
d      4.4
dtype: float64
```

In [6]:

```
# because original1 and original2 have index in different order, Pandas will sort the index before putting them together
df = pd.DataFrame({"x":original1, "y": original2})
df
```

Out[6]:

	x	y
a	3.1	3.1
b	4.2	2.2
c	2.3	1.3
d	1.4	4.4



In [7]:

```
original1.index # the index of original1 is the letters d, c, a, b in a tuple-like object
```

Out[7]:

```
Index(['d', 'c', 'a', 'b'], dtype='object')
```

In [8]:

```
original1['d':'a'] # when slicing pandas uses the index order or original1
```

Out[8]:

```
d      1.4
c      2.3
a      3.1
dtype: float64
```

In [9]:

```
df.index # the index of df are the letters abcd in order
```

Out[9]:

```
Index(['a', 'b', 'c', 'd'], dtype='object')
```

In [10]:

```
df['a':'c'] # when slicing Pandas uses the index order of the DataFrame, which has been sorted
```

Out[10]:

	x	y
a	3.1	3.1
b	4.2	2.2
c	2.3	1.3

## Rearranging value

Both Series and DataFrames have the `.sort_index()` and `.sort_values()` methods which can be used to rearrange the value.

In [11]:

```
original2
```

Out[11]:

```
b      2.2
a      3.1
c      1.3
d      4.4
dtype: float64
```

In [12]:

```
original2.sort_index()
```

Out[12]:



```
a      3.1  
b      2.2  
c      1.3  
d      4.4  
dtype: float64
```

In [13]:

```
original2.sort_values()
```

Out[13]:

```
c      1.3  
b      2.2  
a      3.1  
d      4.4  
dtype: float64
```

In [14]:

```
df
```

Out[14]:

	x	y
a	3.1	3.1
b	4.2	2.2
c	2.3	1.3
d	1.4	4.4

In [15]:

```
df.sort_values(by = "x", ascending = False)
```

Out[15]:

	x	y
b	4.2	2.2
a	3.1	3.1
c	2.3	1.3
d	1.4	4.4

## Changing the Index

The index of a Pandas Series or Pandas DataFrame is immutable and cannot be modified.

However, if you want to change the index of a series or dataframe, you can define a new index and replace the existing index of the series/DataFrame.

In [16]:

```
original1.index = range(4) # I replace the index of the series with this range object.
```

In [17]:

```
original1
```

Out[17]:

```
0      1.4
1      2.3
2      3.1
3      4.2
dtype: float64
```

In [18]:

```
original1.index # We can see this has automatically become a RangeIndex object
```

Out[18]:

```
RangeIndex(start=0, stop=4, step=1)
```

In [19]:

```
original1[1]
```

Out[19]:

2.3

In [20]:

```
original1.loc[1] # behaves the same as above
```

Out[20]:

2.3

In [21]:

```
original1.iloc[1] # behaves the same as above because the range index starts at 0
```

Out[21]:

2.3

In [22]:

```
original1.index = range(1,5)
```

In [23]:

```
original1
```

Out[23]:

```
1      1.4
2      2.3
3      3.1
4      4.2
dtype: float64
```

In [24]:

```
original1[1]
```

Out[24]:

```
1.4
```

In [25]:

```
original1.loc[1]
```

Out[25]:

## 1.4

In [26]:

```
original1.iloc[1] # behavior is different because range index starts at 1
```

Out[26]:

## 2.3

In [27]:

```
original1['a'] # throws an error because 'a' is no longer part of the index and cannot be used to select values
```

-----  
-----

**KeyError**

Traceback

ack (most recent call last)

~\AppData\Local\Temp\ipykernel\_7892\303453168.py

in <module>

----> 1 original1['a'] # throws an error because  
'a' is no longer part of the index and cannot be  
used to select values

~\anaconda3\lib\site-packages\pandas\core\series.py

in \_\_getitem\_\_(self, key)

940

941

elif key\_is\_scalar:

--> 942

return self.\_get\_value(key)



```

943
944         if is_hashable(key):

~\anaconda3\lib\site-packages\pandas\core\series.py in _get_value(self, label, takeable)
1049
1050         # Similar to Index.get_value, but we do not fall back to positional
-> 1051         loc = self.index.get_loc(label)
1052         return self.index._get_values_for_loc(self, loc, label)
1053

~\anaconda3\lib\site-packages\pandas\core\indexes\range.py in get_loc(self, key, method, tolerance)
386         except ValueError as err:
:
387         raise KeyError(key)
from err

```

```
--> 388             raise KeyError(key)
      389         return super().get_loc(key, method=method, tolerance=tolerance)
      390
```

**KeyError:** 'a'

In [28]:

```
original1.index = ['a','b','c','d'] # be careful as no restrictions regarding the meaning of the index is applied.  
# in the original 'a' was associated with 3.1. This index will associate it with 1.4
```

In [29]:

```
original1
```

Out[29]:

```
a      1.4  
b      2.3  
c      3.1  
d      4.2  
dtype: float64
```

In [30]:

```
original1['a']
```

Out[30]:

```
1.4
```

In [31]:

```
original1[0] # now that the index uses strings, you can index by position
```

Out[31]:

**1.4**

In [32]:

```
original1.index = [1, 2, 3, 4, 5] # if the object you provide is of a different length, you get a value error
```

-----  
-----

**ValueError**

Traceback

ack (most recent call last)

~\AppData\Local\Temp\ipykernel\_7892\3051887670.p

y in <module>

----> 1 original1.index = [1, 2, 3, 4, 5] # if the object you provide is of a different length, you get a value error

~\anaconda3\lib\site-packages\pandas\core\generi

c.py in \_\_setattr\_\_(self, name, value)

5498 try:

5499 object.\_\_getattr\_\_(self

, name)

```
-> 5500             return object.__setattr__(self, name, value)
    5501         except AttributeError:
    5502             pass
```

~\anaconda3\lib\site-packages\pandas\\_libs\properties.pyx in pandas.\_libs.properties.AxisProperty.\_\_set\_\_()

```
~\anaconda3\lib\site-packages\pandas\core\series.py in _set_axis(self, axis, labels, fastpath)
    557         if not fastpath:
    558             # The ensure_index call above ensures we have an Index object
--> 559         self._mgr.set_axis(axis, labels)
    560
    561         # ndarray compatibility
```

~\anaconda3\lib\site-packages\pandas\core\intern

```
als\managers.py in set_axis(self, axis, new_labels)
```

```
    214     def set_axis(self, axis: int, new_labels: Index) -> None:
```

```
    215         # Caller is responsible for ensuring we have an Index object.
```

```
--> 216         self._validate_set_axis(axis, new_labels)
```

```
    217         self.axes[axis] = new_labels
```

```
    218
```

```
~\anaconda3\lib\site-packages\pandas\core\internals\base.py in _validate_set_axis(self, axis, new_labels)
```

```
    55
```

```
    56         elif new_len != old_len:
```

```
---> 57             raise ValueError(
```

```
    58                 f"Length mismatch: Expected axis has {old_len} elements, new "
```

```
    59                 f"values have {new_len}
```

```
    59
```

```
    59                 f"values have {new_len}
```

elements"

**ValueError:** Length mismatch: Expected axis has 4 elements, new values have 5 elements



In [33]:

```
# similarly you can change the index of a DataFrame by defining a new object and assigning it to the index.  
df.index = ['j','k','l','m']  
df
```

Out[33]:

	x	y
j	3.1	3.1
k	4.2	2.2
l	2.3	1.3
m	1.4	4.4

# Reindexing

Reindexing is different from just defining a new index.

Reindexing takes a current Pandas object and creates a *new* Pandas object that *conforms* to the specified index.

**Do not confuse reindexing with creating a new index for a dataframe object.**

In [34]:

```
original = pd.Series([1.4, 2.3, 3.1, 4.2], index = ['d', 'c', 'a', 'b'])
```

In [35]:

```
original
```

Out[35]:

```
d      1.4
c      2.3
a      3.1
b      4.2
dtype: float64
```

In [36]:

```
newobj = original.reindex(['a','b','c','d','e']) # note this has an index value that doesn't exist in the original series
```

In [37]:

```
newobj # takes the data in original and moves it so it conforms to the specified index
# values that do not exist for the new index get NaN
```

Out[37]:

```
a      3.1
b      4.2
c      2.3
d      1.4
```

e      NaN

dtype: float64

In [38]:

```
# if you don't want NaN, you can specify a fill_value  
newobj2 = original.reindex(['a','b','c','d','e'], fill_value = 0)  
newobj2
```

Out[38]:

```
a      3.1  
b      4.2  
c      2.3  
d      1.4  
e      0.0  
dtype: float64
```

For ordered data like a time series, it might be desirable to fill values when reindexing

In [39]:

```
obj3 = pd.Series(['blue', 'purple', 'yellow'], index=[0, 3, 6])  
obj3
```

Out[39]:

```
0      blue  
3    purple  
6    yellow  
dtype: object
```

In [40]:

```
obj3.reindex(range(9)) # without any optional arguments, lots of missing values
```

Out[40]:

```
0      blue  
1       NaN  
2       NaN  
3    purple
```

```
4      NaN
5      NaN
6    yellow
7      NaN
8      NaN
dtype: object
```

In [41]:

```
obj3.reindex(range(9), method='ffill')  
# forward-fill pushes values 'forward' until a new value is encountered
```

Out[41]:

```
0      blue  
1      blue  
2      blue  
3    purple  
4    purple  
5    purple  
6    yellow  
7    yellow  
8    yellow  
dtype: object
```

In [42]:

```
obj3.reindex(range(9), method='bfill')  
# back-fill works in the opposite direction  
# there was no value at index 8 so, NaNs get filled in
```

Out[42]:



```
0      blue
1    purple
2    purple
3    purple
4    yellow
5    yellow
6    yellow
7        NaN
8        NaN
dtype: object
```

## Date Ranges as Index

In [43]:

```
# we specify the creation of a date_index using the date_range function
# freq = 'D' creates Daily values
date_index = pd.date_range('1/1/2010', periods=6, freq='D')
date_index
```

Out[43]:

```
DatetimeIndex(['2010-01-01', '2010-01-02', '2010-01-03', '2010-01-04',
               '2010-01-05', '2010-01-06'],
              dtype='datetime64[ns]', freq='D')
```

In [44]:

```
# we create a DataFrame with the date index
df2 = pd.DataFrame({"prices": [100, 101, np.nan, 100, 89, 88]}, index=date_index)
df2
```

Out[44]:

	prices
2010-01-01	100.0
2010-01-02	101.0

prices	
2010-01-03	NaN
2010-01-04	100.0
2010-01-05	89.0
2010-01-06	88.0

In [45]:

```
# we create a DataFrame with the date index
df2 = pd.DataFrame({"prices": [100, 101, np.nan, 100, 89, 88]}, index=date_index)
df2
```

Out[45]:

	prices
2010-01-01	100.0
2010-01-02	101.0
2010-01-03	NaN
2010-01-04	100.0
2010-01-05	89.0
2010-01-06	88.0

In [46]:

```
date_index2 = pd.date_range('12/29/2009', periods=10, freq='D') # a new date index
df2.reindex(date_index2)
```

Out[46]:

	prices
2009-12-29	NaN
2009-12-30	NaN

prices	
2009-12-31	NaN
2010-01-01	100.0
2010-01-02	101.0
2010-01-03	NaN
2010-01-04	100.0
2010-01-05	89.0
2010-01-06	88.0
2010-01-07	NaN

In [47]:

```
df2.reindex(date_index2, method = 'bfill')  
# The value for Jan 3 isn't filled in because that NaN was not created by the reindexing process  
# The NaN already existed in the data.
```

Out[47]:

	prices
2009-12-29	100.0
2009-12-30	100.0
2009-12-31	100.0
2010-01-01	100.0
2010-01-02	101.0
2010-01-03	NaN
2010-01-04	100.0
2010-01-05	89.0
2010-01-06	88.0
2010-01-07	NaN

## .reindex() vs .loc()

If you don't need to fill in any missing info, then `.reindex()` and `.loc()` work very similarly. If the new index will have values that don't exist in the current index, you need to use `reindex`.

In [48]:

```
obj5 = pd.DataFrame({'val':[1.4, 2.3, 3.1, 4.2]}, index = ['d','c','a','b'])  
obj5
```

Out[48]:

	val
d	1.4
c	2.3
a	3.1
b	4.2

In [49]:

```
obj5.reindex(['a','b','c','d'])
```

Out[49]:

	val
a	3.1
b	4.2
c	2.3
d	1.4

In [50]:

```
obj5.loc[['a','b','c','d']] # works the same as reindex
```

Out[50]:

	val
a	3.1
b	4.2
c	2.3
d	1.4

In [51]:

```
obj5.reindex(['a','b','c','d','e'])
```



Out[51]:

	val
a	3.1
b	4.2
c	2.3
d	1.4
e	NaN

In [52]:

```
obj5.loc[['a','b','c','d','e']] # .loc() returns a warning or error if you give an entry in the index that doesn't exist
```

-----  
-----

**KeyError**

Traceback

ack (most recent call last)

~\AppData\Local\Temp\ipykernel\_7892\486090405.py

in <module>

----> 1 obj5.loc[['a','b','c','d','e']] # .loc  
( ) returns a warning or error if you give an entry  
in the index that doesn't exist

~\anaconda3\lib\site-packages\pandas\core\indexing.py in \_\_getitem\_\_(self, key)

929

930 maybe\_callable = com.apply\_i

f\_callable(key, self.obj)

```
--> 931             return self._getitem_axis(ma
ybe_callable, axis=axis)
      932
      933     def _is_scalar_access(self, key: tup
le):
```

```
~\anaconda3\lib\site-packages\pandas\core\indexi
ng.py in _getitem_axis(self, key, axis)
      1151         raise ValueError("Ca
nnot index with multidimensional key")
      1152
-> 1153         return self._getitem_ite
rable(key, axis=axis)
      1154
      1155         # nested tuple slicing
```

```
~\anaconda3\lib\site-packages\pandas\core\indexi
ng.py in _getitem_iterable(self, key, axis)
      1091
      1092         # A collection of keys
```

```
-> 1093         keyarr, indexer = self._get_listlike_indexer(key, axis)
    1094         return self.obj._reindex_with_indexers(
    1095             {axis: [keyarr, indexer]}, copy=True, allow_dups=True
```

```
~\anaconda3\lib\site-packages\pandas\core\indexing.py in _get_listlike_indexer(self, key, axis)
    1312         keyarr, indexer, new_indexer = ax._reindex_non_unique(keyarr)
    1313
-> 1314         self._validate_read_indexer(keyarr, indexer, axis)
    1315
    1316         if needs_i8_conversion(ax.dtype) or isinstance(
```

```
~\anaconda3\lib\site-packages\pandas\core\indexing.py in _validate_read_indexer(self, key, index
```

```
er, axis)
1375
1376         not_found = list(ensure_index(
x(key)[missing_mask.nonzero()[0]].unique())
-> 1377         raise KeyError(f"{not_found}
not in index")
1378
1379
```

**KeyError:** "['e'] not in index"

# Dropping rows or columns

you can use `df.drop()` to remove rows (default) or columns (specify `axis = 1`) at certain index locations.

In [53]:

```
df = pd.DataFrame(np.arange(12).reshape(3,4), columns=['A', 'B', 'C', 'D'], index = ['x','y','z'])
df
```

Out[53]:

	A	B	C	D
x	0	1	2	3
y	4	5	6	7
z	8	9	10	11

In [54]:

```
# drop rows
df.drop(['x', 'z'])
```

Out[54]:

	A	B	C	D
y	4	5	6	7

In [55]:

```
# drop columns  
df.drop(['B', 'C'], axis = 1) # we must specify axis = 1 otherwise Pandas will look for "B" and "C" in the row names
```

Out[55]:

	A	D
x	0	3
y	4	7
z	8	11

In [56]:

```
# df.drop returns a new object and leaves df unchanged  
# you can change this behavior with the argument inplace = True  
df
```

Out[56]:

	A	B	C	D
x	0	1	2	3
y	4	5	6	7
z	8	9	10	11



# Data Alignment

When performing element-wise arithmetic, Pandas will align the index values before doing the computation

In [57]:

```
s1 = pd.Series([7.3, -2.5, 3.4, 1.5], index=['a', 'c', 'd', 'e'])  
s1
```

Out[57]:

```
a      7.3  
c     -2.5  
d      3.4  
e      1.5  
dtype: float64
```

In [58]:

```
s2 = pd.Series([-2.1, 3.6, -1.5, 4, 3.1],  
               index=['a', 'c', 'e', 'f', 'g'])  
s2
```

Out[58]:

```
a    -2.1  
c     3.6  
e    -1.5  
f     4.0  
g     3.1  
dtype: float64
```

In [59]:

```
pd.DataFrame({'s1':s1,'s2':s2}) # for reference
```

Out[59]:

	s1	s2
a	7.3	-2.1
c	-2.5	3.6
d	3.4	NaN
e	1.5	-1.5
f	NaN	4.0
g	NaN	3.1

In [60]:

```
s1 + s2 # returns a new series, where the indexes are the union of the indexes of s1 and s2
```

Out[60]:

a	5.2
c	1.1
d	NaN
e	0.0

```
f      NaN
g      NaN
dtype: float64
```

In [61]:

```
s1.add(s2)
```

Out[61]:

```
a      5.2
c      1.1
d      NaN
e      0.0
f      NaN
g      NaN
dtype: float64
```

In [62]:

```
pd.DataFrame({'s1':s1,'s2':s2})
```

Out[62]:

	s1	s2
a	7.3	-2.1
c	-2.5	3.6
d	3.4	NaN
e	1.5	-1.5
f	NaN	4.0
g	NaN	3.1

In [63]:

```
s1.sub(s2, fill_value = 0)
```

Out[63]:

a	9.4
c	-6.1
d	3.4
e	3.0

```
f    -4.0
g    -3.1
dtype: float64
```

In [64]:

```
s1.rsub(s2, fill_value = 0) # .rsub means 'right hand subtract' sets the series in the argument as the base
```

Out[64]:

```
a    -9.4
c     6.1
d    -3.4
e    -3.0
f     4.0
g     3.1
dtype: float64
```

In [65]:

```
s1 * s2
```

Out[65]:

```
a    -15.33  
c     -9.00  
d         NaN  
e     -2.25  
f         NaN  
g         NaN  
dtype: float64
```

In [66]:

```
s1.multiply(s2, fill_value = 1)
```

Out[66]:

```
a    -15.33  
c     -9.00  
d     3.40  
e     -2.25
```

```
f      4.00  
g      3.10  
dtype: float64
```



For data frames with different columns, the rows and columns will be aligned

In [67]:

```
df1 = pd.DataFrame(np.arange(9.).reshape((3, 3)), columns=list('bcd'),  
                    index=['Ohio', 'Texas', 'Colorado'])  
df1
```

Out[67]:

	b	c	d
Ohio	0.0	1.0	2.0
Texas	3.0	4.0	5.0
Colorado	6.0	7.0	8.0

In [68]:

```
df2 = pd.DataFrame(np.arange(12.).reshape((4, 3)), columns=list('bde'),  
                    index=['Utah', 'Ohio', 'Texas', 'Oregon'])  
df2
```

Out[68]:

	b	d	e
Utah	0.0	1.0	2.0
Ohio	3.0	4.0	5.0
Texas	6.0	7.0	8.0
Oregon	9.0	10.0	11.0



In [69]:

```
df1 + df2
# c is in df1, but not df2
# e is in df2, but not df1
# the result returns the union of columns, but will fill in NaN for elements that do not exist in both
```

Out[69]:

	b	c	d	e
Colorado	NaN	NaN	NaN	NaN
Ohio	3.0	NaN	6.0	NaN
Oregon	NaN	NaN	NaN	NaN
Texas	9.0	NaN	12.0	NaN
Utah	NaN	NaN	NaN	NaN

In [70]:

```
# if you want to fill in values that are missing, you can use df.add() and specify the fill_value
# this will perform the above operation, but instead of using NaN when it can't find a value
# (which will return NaN),
# it will use the fill_value
df1.add(df2, fill_value = 0)
# you still get NaN if the value does not exist in either DataFrame
```

Out[70]:

	b	c	d	e
Colorado	6.0	7.0	8.0	NaN

	b	c	d	e
Ohio	3.0	1.0	6.0	5.0
Oregon	9.0	NaN	10.0	11.0
Texas	9.0	4.0	12.0	8.0
Utah	0.0	NaN	1.0	2.0

Arithmetic operations that can be called on DataFrames and Series are:

- `.add()`, `.radd()` and `.sub()`, `.rsub()`
- `.mul()`, `.rmul()` and `.div()`, `.rdiv()`
- `.floordiv()`, `.rfloordiv()` (floor division `//`)
- `.pow()`, `.rpow()` (exponentiation `**`)

# Summary Stats of a DataFrame

In [71]:

```
df = pd.DataFrame({'one':[1.5,6.0,np.nan, 1.5,4,6, np.nan],  
                  'two':[np.nan, -4.5, np.nan, -1.5, 0, -4.5, 4]},  
                  index=['a', 'b', 'c', 'd','e','f','g'])  
df
```

Out[71]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [72]:

```
df.sum() # default behavior returns column sums and skips missing values  
# default behavior sums across axis 0 (sums the row)
```

Out[72]:

```
one      19.0  
two      -6.5  
dtype: float64
```

In [73]:

```
df # for reference
```

Out[73]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [74]:

```
df.sum(axis = 1) # sum across axis=1, sum across the columns and give row sums
```

Out[74]:

a	1.5
b	1.5
c	0.0



```
d      0.0  
e      4.0  
f      1.5  
g      4.0  
dtype: float64
```

In [75]:

```
df.sum(skipna = False)
```

Out[75]:

```
one     NaN  
two     NaN  
dtype: float64
```

In [76]:

```
df.mean()
```

Out[76]:

```
one    3.8  
two   -1.3  
dtype: float64
```

In [77]:

```
df.mean(axis = 1)
```

Out[77]:

```
a    1.50  
b    0.75  
c    NaN  
d    0.00  
e    2.00  
f    0.75  
g    4.00  
dtype: float64
```



In [78]:

```
df # for reference
```

Out[78]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [79]:

```
df.min()
```

Out[79]:

```
one    1.5
two   -4.5
dtype: float64
```

In [80]:

```
df.idxmin() # which row has the minimum value, also .idxmax()  
# returns the first minimum, if there are multiple  
# you can also specify axis
```

Out[80]:

```
one      a  
two      b  
dtype: object
```

## Summary stats available for dataframes and series

- `count()` - number of non NA values
- `quantile()`
- `sum()`
- `mean()`
- `median()`
- `mad()` - mean absolute deviation
- `prod()`
- `var()`, `std()`

**<https://pandas.pydata.org/pandas-docs/stable/reference/series.html#computations-descriptive-stats>**

## Unique values

In [81]:

```
df # for reference
```

Out[81]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [82]:

```
df.one.unique() # shows the unique values in the order observed
```

Out[82]:

```
array([1.5, 6. , nan, 4. ])
```

In [83]:

```
df.two.unique()
```

Out[83]:

```
array([ nan, -4.5, -1.5,  0. ,  4. ])
```

In [84]:

```
df.unique() # unique can only be applied to a series (a column in a dataframe)
```

-----  
-----

**AttributeError**

Traceback

ack (most recent call last)

~\AppData\Local\Temp\ipykernel\_7892\1052518.py in  
n <module>

----> 1 df.unique() *# unique can only be applied to a series (a column in a dataframe)*

~\anaconda3\lib\site-packages\pandas\core\generic.py in \_\_getattr\_\_(self, name)  
5485 ):



```
5486         return self[name]
-> 5487     return object.__getattr__(s
elf, name)
5488
5489     def __setattr__(self, name: str, val
ue) -> None:
```

**AttributeError:** 'DataFrame' object has no attrib  
ute 'unique'

In [85]:

```
df # for reference
```

Out[85]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [86]:

```
df.one.nunique() # number of non-missing unique values exist
```

Out[86]:

3

In [87]:

```
df.one.value_counts() # tally up counts of each value  
# returns a series. the index are the unique values observed, the values are the frequencies.  
# they appear in descending order of frequency
```

Out[87]:

1.5      2

6.0      2

4.0      1

Name: one, dtype: int64

In [88]:

```
df.one.isin([1.5, 4.0]) # checks to see if the value has membership in a particular list  
# returns a series with boolean values
```

Out[88]:

```
a      True  
b     False  
c     False  
d      True  
e      True  
f     False  
g     False  
Name: one, dtype: bool
```

In [89]:

```
(df.one == 1.5) | (df.one == 4.0) # must use bitwise or. .isin() is much preferred
```

Out[89]:

```
a      True  
b     False  
c     False
```

d      True

e      True

f      False

g      False

Name: one, dtype: bool

In [90]:

```
df.loc[ df.one.isin([1.5,4.0]), ] # can filter rows based on the .isin() membership
```

Out[90]:

	one	two
a	1.5	NaN
d	1.5	-1.5
e	4.0	0.0

## filtering out missing values

In [91]:

```
df
```

Out[91]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [92]:

```
df.dropna() # gets rid of any row that is not complete
```

Out[92]:

	one	two
--	-----	-----

	one	two
b	6.0	-4.5
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5



In [93]:

```
df.dropna(how = 'all') # only drops rows that are entirely NaN
```

Out[93]:

	one	two
a	1.5	NaN
b	6.0	-4.5
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [94]:

```
# you can also use .notnull(), which is True for values that are not missing  
df[df.two.notnull()] # You can use this in conjunction with specifying the column
```

Out[94]:

	one	two
b	6.0	-4.5
d	1.5	-1.5
e	4.0	0.0

	one	two
f	6.0	-4.5
g	NaN	4.0

# Filling in Missing Values

In [95]:

```
df
```

Out[95]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [96]:

```
df.fillna(0) # fill in missing values with a constant
```

Out[96]:

	one	two
--	-----	-----

	one	two
a	1.5	0.0
b	6.0	-4.5
c	0.0	0.0
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	0.0	4.0

In [97]:

```
df.fillna({'one': 1000, 'two': 0}) # use a dictionary to specify values to use for each column
```

Out[97]:

	one	two
a	1.5	0.0
b	6.0	-4.5
c	1000.0	0.0
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	1000.0	4.0

In [98]:

```
df.fillna(method = 'bfill') # backfills. You can also use ffill
```

Out[98]:

	one	two
a	1.5	-4.5
b	6.0	-4.5
c	1.5	-1.5
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [99]:

```
df.mean()
```

Out[99]:

```
one    3.8  
two   -1.3  
dtype: float64
```

In [100]:

```
df.fillna(df.mean()) # fill na with df.mean() will fill in the column means
```

Out[100]:

	one	two
a	1.5	-1.3
b	6.0	-4.5
c	3.8	-1.3
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	3.8	4.0

all of the above fillna methods have created new DataFrame objects. If you want to modify the current DataFrame, you can use the optional argument `inplace = True`



In [101]:

```
df.T
```

Out[101]:

	a	b	c	d	e	f	g
one	1.5	6.0	NaN	1.5	4.0	6.0	NaN
two	NaN	-4.5	NaN	-1.5	0.0	-4.5	4.0

In [102]:

```
# apparently you can only fill missing values with dictionaries/series over a column  
# so we have to do some Transpose magic  
df.T.fillna(df.T.mean()).T
```

Out[102]:

	one	two
a	1.5	1.5
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	4.0	4.0



## dealing with duplicates

In [103]:

```
df
```

Out[103]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
f	6.0	-4.5
g	NaN	4.0

In [104]:

```
df.duplicated() # sees if any of the rows are a duplicate of an earlier row
```

Out[104]:

a	False
b	False
c	False
d	False
e	False
f	True
g	False

dtype: bool

In [105]:

```
df[~df.duplicated()] # gets rid of the duplicated rows
```

Out[105]:

	one	two
a	1.5	NaN
b	6.0	-4.5
c	NaN	NaN
d	1.5	-1.5
e	4.0	0.0
g	NaN	4.0

In [106]:

```
df.one.duplicated()
```

Out[106]:

a	False
b	False
c	False
d	True

e      False

f      True

g      True

Name: one, dtype: bool